

# Processes (ch. 4+5)

Operating Systems

Based on: Three Easy Pieces by Arpaci-Dusseau

Moshe Sulamy

Tel-Aviv Academic College

## A running program

- Lots of processes seemingly running at the same time
- The challenge:
  - Few physical CPUs, illusion of many CPUs

# The Process

- **Virtualizing** the CPU
  - Running one process, stopping it, running another, and so forth
  - **Time sharing** of the CPU
  - Illusion that many virtual CPUs exist

# The Process

- **Virtualizing** the CPU
  - Running one process, stopping it, running another, and so forth
  - **Time sharing** of the CPU
  - Illusion that many virtual CPUs exist
- **Context switch**
  - Low-level mechanism
  - Stop running one program and start running another
- **Scheduling policy**
  - Algorithm to decide which process should run next
  - By history, workload, performance

# Context switch example (xv6-riscv64)

```
1  swtch:
2      sd ra, 0(a0)
3      sd sp, 8(a0)
4      sd s0, 16(a0)
5      sd s1, 24(a0)
6      sd s2, 32(a0)
7      sd s3, 40(a0)
8      sd s4, 48(a0)
9      sd s5, 56(a0)
10     sd s6, 64(a0)
11     sd s7, 72(a0)
12     sd s8, 80(a0)
13     sd s9, 88(a0)
14     sd s10, 96(a0)
15     sd s11, 104(a0)
```

```
1      ld ra, 0(a1)
2      ld sp, 8(a1)
3      ld s0, 16(a1)
4      ld s1, 24(a1)
5      ld s2, 32(a1)
6      ld s3, 40(a1)
7      ld s4, 48(a1)
8      ld s5, 56(a1)
9      ld s6, 64(a1)
10     ld s7, 72(a1)
11     ld s8, 80(a1)
12     ld s9, 88(a1)
13     ld s10, 96(a1)
14     ld s11, 104(a1)
15
16     ret
```

# Time and Space Sharing

- **Time sharing**

- Resource used for a little while by one entity, then a little while by another, and so forth
- e.g., CPU

- **Space sharing**

- Resource is divided (in space) among those who wish to use it
- e.g., memory, disk

# Process vs. Program

- **Program:** static code and static data
- **Process:** dynamic instance of the program
- Multiple processes of the same program can exist

# What constitutes a process?

- Memory (**address space**)
  - Instructions (program code)
  - Data (static and dynamic)
  - `cat /proc/<PID>/maps`
- Registers
  - Program counter (PC)
  - Stack pointer
  - etc.
- I/O information
  - e.g., open files
  - `cat /proc/<PID>/fdinfo/*`




# Process Creation

- Unix like OSes: A process is a replica of a currently existing process.
  - There is a way to load an executable file into an existing process.
- Non-Unix like OSes: A process is created with information from an exe file.

Either way, the first process is created by the OS on initialization.

# Executable Loading



# Executable Loading

- **Load** code and static data into memory
  - Program initially on disk
  - Loading can be done **lazily** (via **paging** and **swapping**)

# Executable Loading

- **Load** code and static data into memory
  - Program initially on disk
  - Loading can be done **lazily** (via **paging** and **swapping**)
- Allocate the **stack**
  - Used for local variables, function parameters, return addresses
  - Initialized with `main` arguments: `argc`, `argv`


# Executable Loading

- **Load** code and static data into memory
  - Program initially on disk
  - Loading can be done **lazily** (via **paging** and **swapping**)
- Allocate the **stack**
  - Used for local variables, function parameters, return addresses
  - Initialized with `main` arguments: `argc`, `argv`
- Allocate the **heap**
  - Used for dynamically-allocated data
  - Request space by calling `malloc`, free it by `free`

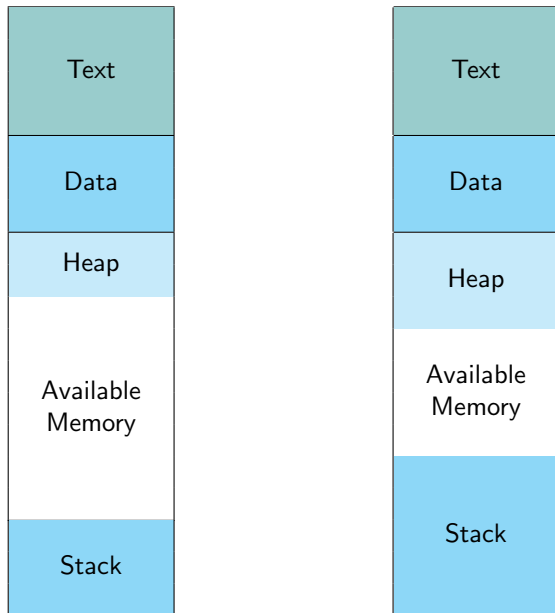
# Executable Loading

- **Load** code and static data into memory
  - Program initially on disk
  - Loading can be done **lazily** (via **paging** and **swapping**)
- Allocate the **stack**
  - Used for local variables, function parameters, return addresses
  - Initialized with `main` arguments: `argc`, `argv`
- Allocate the **heap**
  - Used for dynamically-allocated data
  - Request space by calling `malloc`, free it by `free`
- Start program at entry point (NOT necessarily `main()`)
  - Transfer control of CPU to the newly-created process

# Executable loading




# Two Processes for One Program (Virt Mem View)

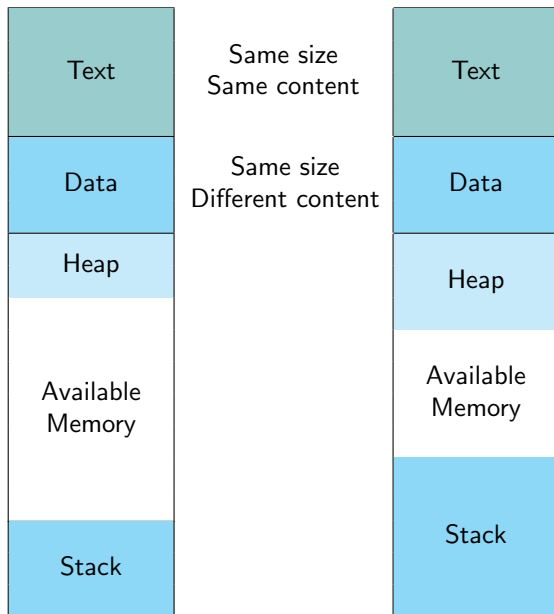




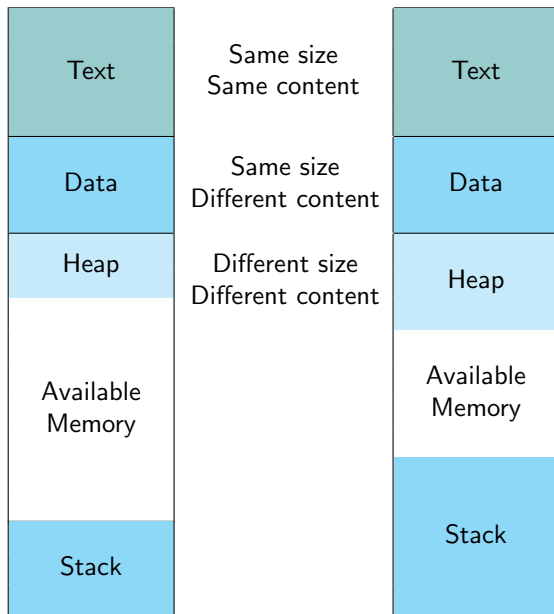
# Two Processes for One Program (Virt Mem View)



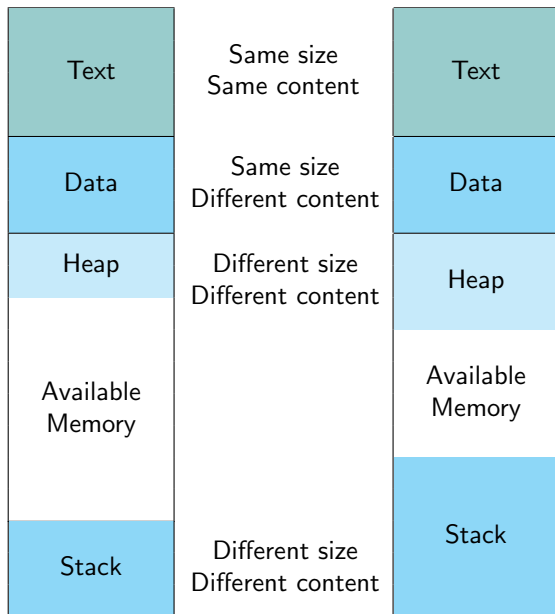
# Two Processes for One Program (Virt Mem View)



# Two Processes for One Program (Virt Mem View)



# Two Processes for One Program (Virt Mem View)



# Process Life Cycle

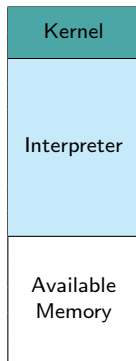
- A process can be in a finite number of **states**
- Events cause **transitions** between states

# Process Life Cycle

- A process can be in a finite number of **states**
- Events cause **transitions** between states
- **Single-tasking** OS (Ancient, Small phys. memory)
  - Only one process at a time
  - **Interpreter** loaded on boot, overwrites part of itself into process


# Process Life Cycle

- A process can be in a finite number of **states**
- Events cause **transitions** between states
- **Single-tasking** OS (Ancient, Small phys. memory)
  - Only one process at a time
  - **Interpreter** loaded on boot, overwrites part of itself into process



# Process Life Cycle

- A process can be in a finite number of **states**
- Events cause **transitions** between states
- **Single-tasking** OS (Ancient, Small phys. memory)
  - Only one process at a time
  - **Interpreter** loaded on boot, overwrites part of itself into process

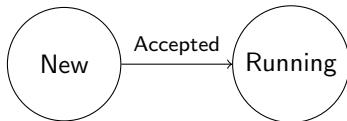




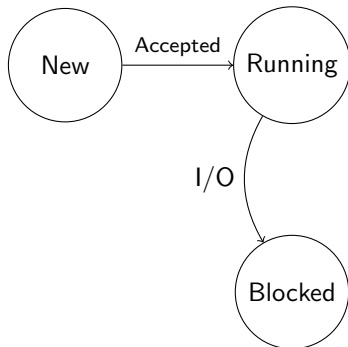
# Process Life Cycle - Single Tasking




# Process Life Cycle - Single Tasking



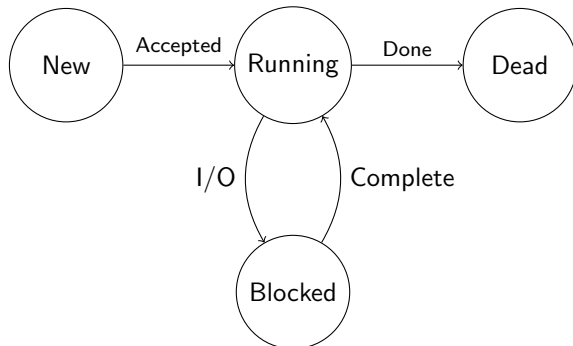
# Process Life Cycle - Single Tasking



# Process Life Cycle - Single Tasking



# Process Life Cycle - Single Tasking



# Process Life Cycle


- Modern operating systems: **multi-tasking**
  - Multiple processes co-exist
  - **Cooperative** multi-tasking: `yield`
  - **Preemptive** multi-tasking: **interrupts**

# Process Life Cycle

- Modern operating systems: **multi-tasking**
  - Multiple processes co-exist
  - **Cooperative** multi-tasking: `yield`
  - **Preemptive** multi-tasking: **interrupts**
- A process can be **ready** to run, but not running
  - OS schedules a process to run for a while, then deschedules it and picks another process, and so forth
  - A new state: **ready**

# Process States

- **Running:** executing on CPU
- **Ready:** ready to run, waiting to be scheduled
- **Blocked:** suspended, waiting for some event





# Process States - Example I

Time	Process 0	Process 1	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	
4	Running	Ready	Process 0 done
5	-	Running	
6	-	Running	
7	-	Running	
8	-	Running	Process 1 done

# Process States - Example II

Time	Process 0	Process 1	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	0 initiates I/O
4	Blocked	Running	0 is blocked
5	Blocked	Running	so 1 runs
6	Blocked	Running	
7	Ready	Running	I/O done
8	Ready	Running	Process 1 done
9	Running	-	
10	Running	-	Process 0 done

# Data Structures

- OS maintains a data structure of active processes
  - The **process table**
  - Limited size - `cat /proc/sys/kernel/threads-max`

# Data Structures

- OS maintains a data structure of active processes
  - The **process table**
  - Limited size - `cat /proc/sys/kernel/threads-max`
- Process Control Block (**PCB**):
  - Process identifier (**PID**)
  - State
  - Related processes (parent)
  - CPU context, e.g., registers (saved when suspended)
  - Memory locations
  - Open files

# Summary (Process Abstraction)

- **Process:** OS abstraction of a running program
- Can be described by:
  - **Address space**
  - CPU registers (inc. **program counter** & **stack pointer**)
  - I/O information (e.g., open files)
- **Process state:** running, ready to run, blocked.
  - transition by different events
- **Process list:** information about all processes in the system
  - **Process control block:** a structure with information about a specific process

# Process API

- API: Application Programming Interface
- The API of the OS: **system calls**
  - Function call into OS code
  - Higher privilege level, for sensitive operations (e.g., hardware)

# Process API

- API: Application Programming Interface
- The API of the OS: **system calls**
  - Function call into OS code
  - Higher privilege level, for sensitive operations (e.g., hardware)
- Rewrite code for each OS?
  - **POSIX API**: standard set for each POSIX-compliant OS  
`write)`

# POSIX hides details

## fork xv6-x86

```
1 movl    $1, %eax
2 int     $64
```

## fork Linux-x86

```
1 movl    $2, %eax
2 int     $128
```

## close xv6-x86

```
1      pushl    fd
2      subl     $4,%esp
3      movl     $21,%eax
4      int      $64
5      addl     $4,%esp
```

## close Linux-x86

```
1      movl     fd,%ebx
2      movl     $6,%eax
3      int      $127
```



# Posix Process API

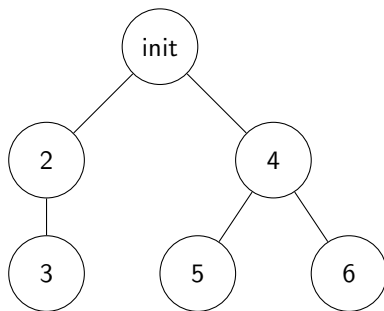
- `fork()`: create a new process
- `wait()`: block until a child process terminates
- `exec()`: make the process execute a given program

# Process Tree

- Start with one process: `init` (PID 1)
- A process can create processes
  - Process *A* creates *B*: *A* is the **parent** of *B*, *B* is the **child** of *A*
  - Can create many children, only one parent
  - Parent can **wait** for child process to finish
- Process ID (**PID**): increasing identifier
  - Get PID: `getpid()`
  - Get parent PID: `getppid()`

# Process Tree

- Processes form a tree:



- `ps --forest -eaf`
- `pstree`

# fork()

- `fork()`: creates a new process
  - Wrapper for `clone` (in Linux)
- New process: almost exact copy of parent
  - Same: memory, execution point, open files
  - Different: PID, return value
  - **Copy-on-write** (Optimization)

# fork()

- `fork()`: creates a new process
  - Wrapper for `clone` (in Linux)
- New process: almost exact copy of parent
  - Same: memory, execution point, open files
  - Different: PID, return value
  - **Copy-on-write** (Optimization)
- Parent: `fork()` returns an integer:
  - If successful returns the **PID** of created child process
  - If fails negative number for error code

# fork()

- `fork()`: creates a new process
  - Wrapper for `clone` (in Linux)
- New process: almost exact copy of parent
  - Same: memory, execution point, open files
  - Different: PID, return value
  - **Copy-on-write** (Optimization)
- Parent: `fork()` returns an integer:
  - If successful returns the **PID** of created child process
  - If fails negative number for error code
- Child process:
  - Begins to run at the point after the fork.
  - 'return value' is zero.

# fork in details

1 `pid = fork();`

1 `movl $1,%eax`  
2 `int $64`  
3 `movl %eax,pid`

Parent		Child	
1	<code>movl \$1,%eax</code>		
2	<code>int \$64</code>		
1	<code>movl %eax,pid</code>	1	<code>movl %eax,pid</code>

# fork()

## Typical usage example (fork.c):

```
1 printf("hello world (pid:%d)\n", getpid());
2 int rc = fork();
3 if (rc < 0) {
4     fprintf(stderr, "fork failed\n");
5     exit(1);
6 }
7 else if (rc == 0) {
8     // child (new process)
9     // sleep(5);    // Try with and without
10    printf("I am child of %d (pid:%d)\n", getppid(), getpid());
11 }
12 else {
13     // parent
14    printf("I am parent of %d (pid:%d)\n", rc, getpid());
15 }
```



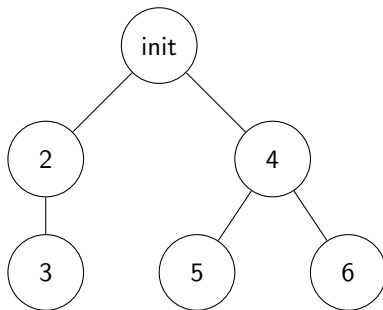
# fork()

## Output:

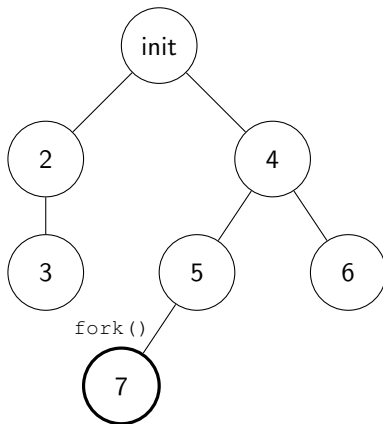
```
prompt> gcc -o fork fork.c -Wall
prompt> ./fork
hello world (pid:1300)
I am parent of 1301 (pid:1300)
I am child of 1 (pid:1301)
prompt>
```

- Child of **1??**


# fork()



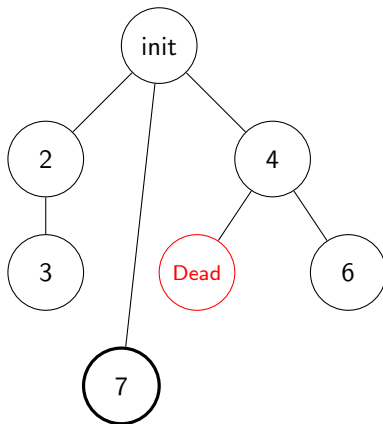
# fork()



# fork()



# fork()



# fork()

peculiar1.c:

```
1  int main(int argc, char *argv[])
2  {
3      fork();
4      fork();
5      printf("hello there\n");
6  }
```

What is the output?

# fork()

peculiar1.c:

```
1  int main(int argc, char *argv[])
2  {
3      fork();
4      fork();
5      printf("hello there\n");
6  }
```

What is the output?

```
1  hello there
2  hello there
3  hello there
4  hello there
```

# fork()

peculiar2.c:

```
1  int main(int argc, char *argv[])
2  {
3      int pid = fork();
4      if (pid)
5          fork();
6      fork();
7      printf("hello there\n");
8  }
```

What is the output?



# fork()

peculiar2.c:

```
1  int main(int argc, char *argv[])
2  {
3      int pid = fork();
4      if (pid)
5          fork();
6      fork();
7      printf("hello there\n");
8  }
```

What is the output?


P

# fork()

peculiar2.c:

```
1  int main(int argc, char *argv[])
2  {
3      int pid = fork();
4      if (pid)
5          fork();
6      fork();
7      printf("hello there\n");
8  }
```

What is the output?




# fork()

peculiar2.c:

```
1  int main(int argc, char *argv[])
2  {
3      int pid = fork();
4      if (pid)
5          fork();
6      fork();
7      printf("hello there\n");
8  }
```

What is the output?




# fork()

peculiar2.c:

```
1  int main(int argc, char *argv[])
2  {
3      int pid = fork();
4      if (pid)
5          fork();
6      fork();
7      printf("hello there\n");
8  }
```

What is the output?



# fork()

peculiar3.c:

```
1  int main(int argc, char *argv[])
2  {
3      fork();
4      printf("hello\n");
5  }
```

Can this print "hehellollo"?

# fork()

peculiar3.c:

```
1  int main(int argc, char *argv[])
2  {
3      fork();
4      printf("hello\n");
5  }
```

Can this print "hehellollo"?

- No! Due to how `printf` works
- But... very important to consider these cases
- More on this in the future (**concurrency**)

# fork()

peculiar4.c:

```
1  int main(int argc, char *argv[])
2  {
3      int x = 0;
4      if (fork()) {
5          sleep(5); // BLOCKED state for 5 seconds
6          printf("%d\n", x);
7      }
8      else {
9          x += 3;
10     }
11 }
```

What is the output?

# fork()

peculiar4.c:

```
1  int main(int argc, char *argv[])
2  {
3      int x = 0;
4      if (fork()) {
5          sleep(5); // BLOCKED state for 5 seconds
6          printf("%d\n", x);
7      }
8      else {
9          x += 3;
10     }
11 }
```

What is the output? **0**

- Why?



# fork()

peculiar4.c:

```
1  int main(int argc, char *argv[])
2  {
3      int x = 0;
4      if (fork()) {
5          sleep(5); // BLOCKED state for 5 seconds
6          printf("%d\n", x);
7      }
8      else {
9          x += 3;
10     }
11 }
```

What is the output? **0**

- Why? Child's memory is a **copy**

# fork()

peculiar5.c:

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

How many processes does this code create?

# fork()

peculiar5.c:

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

How many processes does this code create?

P

# fork()

peculiar5.c:

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

How many processes does this code create?




# fork()

peculiar5.c:

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

How many processes does this code create?




# fork()

peculiar5.c:

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

How many processes does this code create?




# fork()

peculiar5.c:

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

How many processes does this code create?



# fork()

peculiar6.c:

```
1  int main(int argc, char *argv[])
2  {
3      int x = 0;
4      if (fork()) {
5          sleep(5); // BLOCKED state for 5 seconds
6      }
7      else {
8          x += 3;
9      }
10     printf("%d\n", x);
11 }
```

Last one - what is the output?



# fork()

peculiar6.c:

```
1  int main(int argc, char *argv[])
2  {
3      int x = 0;
4      if (fork()) {
5          sleep(5); // BLOCKED state for 5 seconds
6      }
7      else {
8          x += 3;
9      }
10     printf("%d\n", x);
11 }
```

Last one - what is the output? **30** or **03**

- Depends on scheduling

# fork()

peculiar6.c:

```
1  int main(int argc, char *argv[])
2  {
3      int x = 0;
4      if (fork()) {
5          sleep(5); // BLOCKED state for 5 seconds
6      }
7      else {
8          x += 3;
9      }
10     printf("%d\n", x);
11 }
```

Last one - what is the output? **30** or **03**

- Depends on scheduling
- Can we make it deterministic?

# wait()

- `wait()`: waits for a child process to finish
  - Any child process (if several exist)
  - Returns PID of terminated child process (-1 on error)
  - `waitpid()`: waits for a specific child process (by PID)
- To wait for all child processes to end:
  - `while (wait(NULL) != -1);`

# wait()

wait.c:

```
1  int main(int argc, char *argv[])
2  {
3      int x = 0;
4      int rc = fork();
5      if (rc) {
6          wait(NULL); // BLOCKED until child terminates
7          // equivalent here: waitpid(rc, NULL, 0);
8      }
9      else {
10         x += 3;
11     }
12     printf("%d\n", x);
13 }
```

Output is always **30**

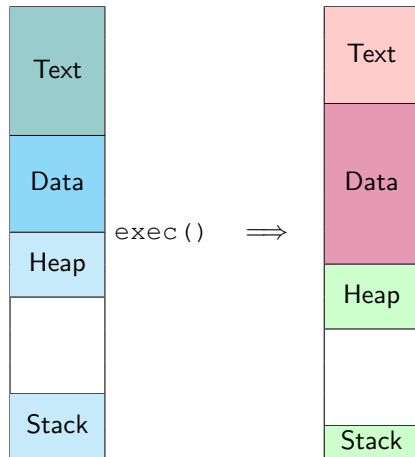
# exec()

- After `fork()`, parent and child execute same code
  - What if we want to run a different program?
  - `exec()` does just that
- Six variants of `exec()`: `execl`, `execlp`, `execle`, `execv`, `execvp`, `execvpe`. Read man for details

# exec()

- After `fork()`, parent and child execute same code
  - What if we want to run a different program?
  - `exec()` does just that
- Six variants of `exec()`: `execl`, `execlp`, `execle`, `execv`, `execvp`, `execvpe`. Read man for details
- `exec()`: transform current program into a different program
  - Receives program name and arguments (`argv`)
  - Overwrites and re-initializes process memory
  - A successful `exec()` never returns!

# exec()



# exec()

exec.c:

```
1  int main(int argc, char *argv[])
2  {
3      int rc = fork();
4      if (rc < 0) {
5          fprintf(stderr, "fork failed\n");
6          exit(1);
7      }
8      else if (rc == 0) {
9          char* args[4] = { "wc", "-l", "exec.c", NULL };
10         execvp(args[0], args);
11         printf("this shouldn't print out\n");
12     }
13     else {
14         int rc_wait = wait(NULL); // or waitpid(rc, NULL, 0)
15         printf("I am parent of %d (rc_wait:%d) (pid:%d)\n",
16             rc, rc_wait, getpid());
17     }
18 }
```



# The Living Dead

- When a process terminates, it remains in the process list as a **zombie**
  - Parent process may want to know its status
- Zombie remains until it is reaped (or its parent terminates)
  - Process 1 adpots orphans (zombied or live)
- A program should not leave zombies!



# The Living Dead

- How to avoid zombies?
  - `wait()`: blocks until a child completes & reaps it
  - `waitpid()`: blocks until a specific child completes & reaps it
- Not enough
  - The terminal (shell) executes processes in the background, wants to continue accepting user input
  - It is possible to `wait()` without blocking, but very inconvenient
- What can we do?



- **Software interrupts**

- Asynchronous notification of an event
- Inter-process communication (**IPC**) or messages from OS

- **Software interrupts**

- Asynchronous notification of an event
- Inter-process communication (**IPC**) or messages from OS

- Various signals exist:

- ^C in the terminal sends `SIGINT` (“interrupt from keyboard”)
- Invalid memory reference causes `SIGSEGV`
- A process can send `SIGKILL` to another process
- Child process terminated - **SIGCHLD**

# Signal Handlers

- Some signals are handled automatically by the OS
  - `SIGKILL`, `SIGSTOP`
- Others are handled by a **signal handler**
  - Each signal has a default behavior, e.g., `SIGINT` causes the process to terminate
  - Can override default with `sigaction()`
- Let's write our own **signal handler**!

# Signal Handlers

signal1.c:

```
1  int main(int argc, char *argv[])
2  {
3      struct sigaction act;
4      sigemptyset(&act.sa_mask);
5      act.sa_handler = SIG_IGN;
6      act.sa_flags = 0;
7
8      if (sigaction(SIGINT, &act, NULL) == -1) {
9          fprintf(stderr, "sigaction failed\n");
10         exit(1);
11     }
12     while (1);
13 }
```

# Signal Handlers

signal2.c:

```
1 void signal_handler(int signal) {
2     if (signal == SIGCHLD) {
3         int rc = wait(NULL);
4         printf("child terminated %d (pid:%d)\n", rc, getpid());
5     }
6 }
7 int main(int argc, char *argv[])
8 {
9     struct sigaction act;
10    sigemptyset (&act.sa_mask);
11    act.sa_handler = signal_handler;
12    act.sa_flags = 0;
13
14    sigaction(SIGCHLD, &act, NULL);
15    if (fork()) {
16        while (1);
17    }
18 }
```

No zombies!

# kill()

- `kill()`: send a signal to another process
  - `kill(pid_t pid, int sig)`
  - *pid*: process id to send signal to
  - *sig*: signal to send
- Name is misleading
  - Can send any signal



# Case Study

- How does a shell work?
  - Reads user command
  - Forks a child
  - Sets up process (e.g., redirection)
  - Execs the relevant program
  - Waits for it to finish (if not background)
  - Reads next command

# Summary (Process API)

- `fork()`: create a new process (clone current)
- `wait()`: waits for a child process to finish
  - Also `waitpid()`
- `exec()`: transform program into a different program
  - Successful `exec()` never returns
- Terminated process remains as a **zombie**, to avoid:
  - Parent terminates
  - `wait()` or `waitpid()` by parent
- **Signals** are software interrupts
  - Can write our own **signal handlers**
  - Also helps with zombies
- `kill()`: send a signal to another process