

# Virtual Memory

Operating Systems


Based on: Three Easy Pieces by Arpaci-Dusseau

Moshe Sulamy

Tel-Aviv Academic College


# Processes & Memory

- Each process has its own **address space**
  - 4GB (32-bit)
    - Not really. Linux takes top 1GB. Windows top 2GB.
  - Heap, Stack, Data, Code (Text)




# Processes & Memory

- Each process has its own **address space**
  - 4GB (32-bit)
    - Not really. Linux takes top 1GB. Windows top 2GB.
  - Heap, Stack, Data, Code (Text)
- Heap
  - Dynamic memory allocation
- Stack
  - Automatic memory allocation



# Processes & Memory

- Each process has its own **address space**
  - 4GB (32-bit)
    - Not really. Linux takes top 1GB. Windows top 2GB.
  - Heap, Stack, Data, Code (Text)
- Heap
  - Dynamic memory allocation
- Stack
  - Automatic memory allocation
- Data
  - Static (Global/Local) values and variables
- Code
  - Program instructions



# Processes & Memory

- Each process has its own **address space**
  - 4GB (32-bit)
- Ten processes: 40GB of RAM!
  - Typically we have much less physical memory
  - And many more processes

How can the OS provide a private, potentially large address space for multiple running processes?

# Virtualizing Memory

- OS virtualizes physical memory
- Goals:
  - **Transparency:**
    - Invisible to the running program
    - A process “thinks” it has a continuous **address space**
  - **Efficiency:**
    - Not making programs run much more slowly
    - Not using too much memory to support virtualization
  - **Protection:**
    - Protect processes from one another, and the OS from processes
    - **Isolation** among processes

# Address Translation

- **Hardware-based address translation**

- On every memory reference, address translation is performed
- Hardware redirects memory references to physical locations

- OS manages memory locations

- Which are free and which are in use

- Hardware support

- e.g., registers, TLBs, page-table


# Assumptions

- ① Address space must be placed contiguously
- ② The size is less than the physical memory size
- ③ Each address space is the same size



# Dynamic Relocation

- Also called: **base and bounds**
  - Hardware registers: **base** and **bounds**
  - OS decides where in physical memory a process is loaded
    - Sets **base** register to that value
  - Memory references are **translated** by **base**
    - $\text{physical address} = \text{virtual address} + \text{base}$
  - Processor checks reference is within **bounds** of base




# Hardware Support

- Two (or more) CPU modes:
  - OS runs in **privileged mode** (or **kernel mode**)
  - Applications run in **user mode**
- **Base** and **bounds** registers
  - Hardware is called **memory management unit (MMU)**
- Generate **exceptions** on illegal access
  - Execute OS **exception handler**

- A process starts running:
  - Find space for address space in physical memory
  - Maintain a **free list** of free address spaces
- A process is terminated:
  - Reclaim the memory for use
  - Add it back to the **free list**
- Context switch:
  - Save and restore the base and bounds registers
  - In the process structure (PCB)



# Segmentation

- Base and bounds is problematic
  - Big chunk of free space
  - Still taking up physical memory
- Solution?



# Segmentation

- Base and bounds is problematic
  - Big chunk of free space
  - Still taking up physical memory
- Solution? **segmentation**
  - Base and bounds pair for each segment



# Segmentation

- A contiguous portion of the address space
  - Logical segments: code, stack, heap
- Each can be placed in a different part of the physical memory
- $\text{physical address} = \text{offset} + \text{base}$ 
  - Not  $\text{virtual address} + \text{base}$ !
  - e.g., offset of virtual address 100 is 100
  - Offset of virtual address 4200 can be 104
    - Since it is 104 in the heap segment

# Segmentation

- Ever encountered a **segmentation fault**?
- If an **illegal address** beyond the segment is referenced:
  - Hardware detects **out of bounds** access
  - OS event: **segmentation fault**

# Support for Sharing

- Segment can be shared between address spaces (processes)
  - **Code sharing** is common
  - Same program, no need to load the code twice
- Extra hardware support: **protection bits**
  - Code segment is read-only
  - Can be shared without harming **isolation**

Segment	Base	Size	Grows	Positive?	Protection
Code	32K	2K		1	Read-Execute
Heap	34K	3K		1	Read-Write
Stack	28K	2K		0	Read-Write



# Fine-Grained vs. Coarse-Grained

- Thus far: just a few segments
  - Code, stack, heap
  - **Coarse-grained**
- **Fine-grained**
  - Large number of smaller segments
  - **Segment table** stored in memory
  - More flexible

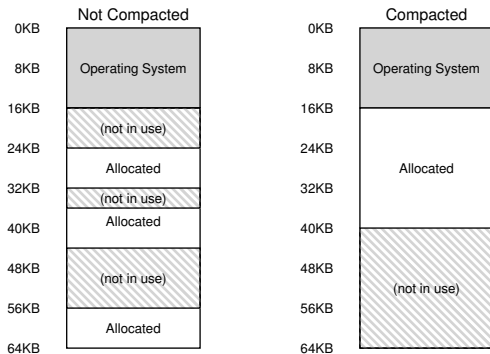
- On context switch: segment registers saved and restored
- On memory allocation/free: update segment size

- On context switch: segment registers saved and restored
- On memory allocation/free: update segment size
- **External fragmentation**
  - Physical memory becomes full of little holes of free space
  - Difficult to allocate new segments, or grow existing ones

# Compaction


- **Compaction**

- Stop running processes
- Copy data to contiguous region
- Change segment registers accordingly
- Compaction is expensive!



# External Fragmentation

- Detour to discuss **free-space management**
- Also applies to user-level memory allocation
  - e.g., `malloc()` and `free()`
- Not a problem with fixed-size chunks
  - Can use a free-list



# External Fragmentation

Assume:

- A basic heap interface:
  - `malloc(size_t size)` allocates `size` or more bytes
  - `free(void *ptr)` frees corresponding chunk
    - Note that no size is provided
- Only **external fragmentation**
  - Allocators also have **internal fragmentation**
  - Unused space in chunks bigger than requested
- No memory relocation
  - No **compaction** of free space

# Splitting



# Splitting

`malloc(2048)`





# Splitting

`malloc(2048)`

split



# Splitting

```
malloc(2048)=15KB
```



# Splitting



# Coalescing

free (15KB)



# Coalescing

free (15KB)



# Coalescing

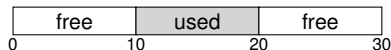


# Coalescing




# Splitting

- Find a free chunk to satisfy request and split it into two
  - Assume the following 30-byte heap:



- Its free list is:

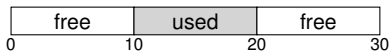


- After a 1-byte request:




# Splitting


- Find a free chunk to satisfy request and split it into two
  - Assume the following 30-byte heap:



- Its free list is:

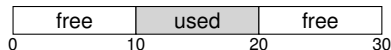


- After a 1-byte request:



# Coalescing

- Coalesce free space when memory is freed
  - i.e., merge contiguous free chunk
  - Consider our previous heap:



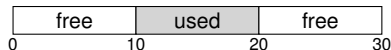
- After a call to `free(10)`:



- With coalescing:

# Coalescing

- Coalesce free space when memory is freed
  - i.e., merge contiguous free chunk
  - Consider our previous heap:



- After a call to `free(10)`:





- With coalescing:



# Tracking The Size

- Interface to `free(void *ptr)` does not provide size
- Store extra information in a **header** block
  - Usually just before chunk of memory
  - Magic number for integrity checking




# Free List

- Need to implement the **free list** itself
  - Can't call `malloc()` - we are implementing it!
  - Need to embed the list inside the free space itself


# Free List

- Need to implement the **free list** itself
  - Can't call `malloc()` - we are implementing it!
  - Need to embed the list inside the free space itself
- Example: manage 4096-byte chunk
  - Maintain `size` and `next` for each node:



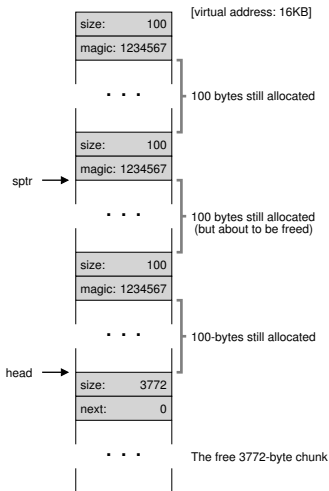
# Free List

- 100 bytes are requested: **split** chunk



# Free List


- Three allocated regions of 100 bytes:






# Free List

- After `free(16500)`:
  - Region start: 16384, previous chunk: 108, current header: 8



# Free List

- Free last two chunks: fragmentation!
  - Need to **coalesce** the list



# Managing Free Space

- **Best Fit**

- Return smallest chunk that's as big or bigger than requested size
- Exhaustive search: heavy performance penalty

- **Worst Fit**

- Opposite of best fit: return largest chunk
- Still requires full search, bad performance, excess fragmentation

# Managing Free Space

- **First Fit**

- Return first block that is big enough
- Speed advantage, but pollutes beginning of list

- **Next Fit**

- As first fit, but start where stopped previously
- Spreads the searches throughout the list

# Managing Free Space

- Examples: allocation request size 15



- Best-fit:



- Worst-fit:




- First-fit: same as worst-fit, but faster

# Segregated Lists

- Keep lists for fixed-size objects
- General memory allocator for the rest
- **Slab allocator** allocates **object caches**
  - For common kernel objects (locks, file-system inodes, etc.)
  - When a cache is running low: request **slab** of memory from general allocator

# Buddy Allocation

- Make coalescing simple: **binary buddy allocator**
- Divide free space by two until a block is found
  - Further split into two is too small
  - Suffers from **internal fragmentation**
  - Easy to coalesce:
    - Recursively up the tree
    - Buddy address differs by a single bit



# Summary

- Virtualize RAM into process **address space**
- Address translation:
  - Dynamic relocation (**base and bounds**)
  - Segmentation
    - **Coarse-grained**: just a few segments
    - **Fine-grained**: large number of smaller segments
- **External fragmentation**
  - Free memory fragments into small parts
  - Splitting & coalescing
- **Compaction**
  - Stop processes, copy data to contiguous region