

Virtual Memory

Operating Systems

Based on: Three Easy Pieces by Arpaci-Dusseau

Moshe Sulamy

Tel-Aviv Academic College

Processes & Physical Addresses


Can physical addresses be used while running? Yes. However:

- One process in memory at a time:
 - OS loads process to a fixed address? Cumbersome
 - OS loads at arbitrary place? Position independent code
- Several process in memory at a time:
 - Code must be position independent
 - Process can access other process memory
 - External fragmentation: Process cannot load.
 - Must wait for processes to terminate

Physical addresses and 'process' addresses need to be separate

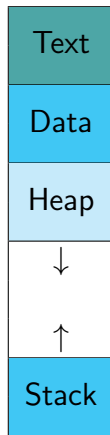
Process View of Memory

- 4GB (32-bit)
 - Not really. Linux takes top 1GB. Windows top 2GB.
- Heap, Stack, Data, Code (Text)



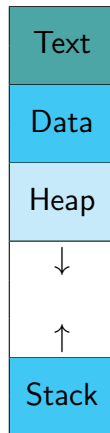
Process View of Memory

- 4GB (32-bit)
 - Not really. Linux takes top 1GB. Windows top 2GB.
- Heap, Stack, Data, Code (Text)
- Heap
 - Dynamic memory allocation
- Stack
 - Automatic memory allocation



Process View of Memory

- 4GB (32-bit)
 - Not really. Linux takes top 1GB. Windows top 2GB.
- Heap, Stack, Data, Code (Text)
- Heap
 - Dynamic memory allocation
- Stack
 - Automatic memory allocation
- Data
 - Static (Global/Local) values and variables
- Code
 - Program instructions



Processes vs. Memory

- Each process has its own **address space**
 - 4GB (32-bit)
- Ten processes: 40GB of RAM!
 - Typically we have much less physical memory
 - And many more processes

How can the OS provide a private, potentially large address space for multiple running processes?

Virtualizing Memory

- OS and hardware virtualizes physical memory
- Goals:
 - **Transparency:**
 - Invisible to the running program
 - A process “thinks” it has a continuous **address space**
 - **Efficiency:**
 - Not making programs run much more slowly
 - Not using too much memory to support virtualization
 - **Protection:**
 - Protect processes from one another, and the OS from processes
 - **Isolation** among processes

Address Translation

- **Hardware-based address translation**

- On every memory reference, address translation is performed
- Hardware redirects memory references to physical locations

- OS manages memory locations

- Which are free and which are in use

- Hardware support

- e.g., registers, TLBs, page-table

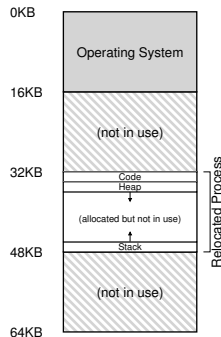
Assumptions

- ① Address space must be placed contiguously
- ② The size is less than the physical memory size
- ③ Each address space is the same size

Dynamic Relocation

Also called: **base and bounds**

- Hardware registers: **base** and **bounds**
- OS decides where in physical memory a process is loaded
 - Sets **base** register to that value
 - Sets **bounds** register
- Memory references are **translated** by MMU
 - $\text{physical address} = \text{virtual address} + \text{base}$
- MMU checks reference is within **bounds** of base



Hardware Support

- Two (or more) CPU modes:
 - OS runs in **privileged mode** (or **kernel mode**)
 - Applications run in **user mode**
- **Base** and **bounds** registers
 - Hardware is called **memory management unit (MMU)**
- Generate **exceptions** on illegal access
 - Execute OS **exception handler**

We are in better situation

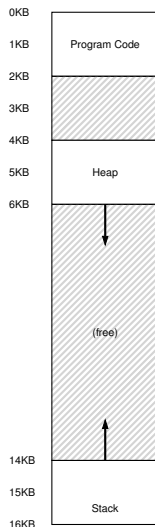
- Solved cross process memory access
- External fragmentation exists but situation is better
 - compactification (not reasonable for huge memories)
- Code can believe it is loaded in the same address always
- Simple and efficient.

Same issues as before relocation

- A process starts running:
 - Find space for address space in physical memory
 - Maintain a **free list** of free address spaces
- A process is terminated:
 - Reclaim the memory for use
 - Add it back to the **free list**

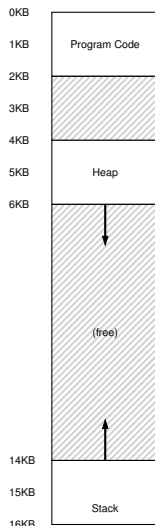

Segmentation

- Base and bounds is problematic
 - Big chunk of free space
 - Still taking up continuous physical memory
- Solution?



Segmentation

- Base and bounds is problematic
 - Big chunk of free space
 - Still taking up continuous physical memory
- Solution? **segmentation**
 - Base and bounds pair for each segment



Segmentation

- A contiguous portion of the address space
 - Logical segments: code, stack, heap
- In essence each segment has its own address space
- Each can be placed in a different part of the physical memory
- Usually under program (or rather compiler) control
- $\text{physical address} = \text{offset} + \text{base}$
 - Not $\text{virtual address} + \text{base}$!
 - e.g., offset of virtual address 100 is 100
 - Offset of virtual address 4200 can be 104
 - Since it is 104 in the heap segment

Segmentation

- Ever encountered a **segmentation fault**?
- If an **illegal address** beyond the segment is referenced:
 - Hardware detects **out of bounds** access
 - OS event: **segmentation fault**

Support for Sharing

- Segment can be shared between address spaces (processes)
 - **Code sharing** is common
 - Same program, no need to load the code twice
- Extra hardware support: **protection bits**
 - Code segment is read-only
 - Can be shared without harming **isolation**

Segment	Base	Size	Grows	Positive?	Protection
Code	32K	2K		1	Read-Execute
Heap	34K	3K		1	Read-Write
Stack	28K	2K		0	Read-Write

Fine-Grained vs. Coarse-Grained

- Thus far: just a few segments
 - Code, stack, heap
 - **Coarse-grained**
- **Fine-grained**
 - Large number of smaller segments
 - **Segment table** stored in memory
 - More flexible

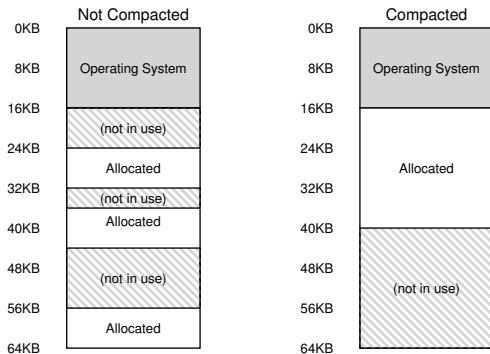
- On memory allocation/free: update segment size

- On memory allocation/free: update segment size
- **External fragmentation**
 - Physical memory becomes full of little holes of free space
 - Difficult to allocate new segments, or grow existing ones
 - But possible: Compactification

Compaction

- **Compaction**

- Stop running processes
- Copy used memory to contiguous region
- Change segment registers accordingly
- Compaction is expensive!




Dynamic relocation and
segmentation are relic of the
past

Detour - Heap management

External Fragmentation

- Detour to discuss **free-space management**
- Also applies to user-level memory allocation
 - e.g., `malloc()` and `free()`
- Not a problem with fixed-size chunks
 - Can use a free-list



20 free bytes fragmented into chunks of 10

External Fragmentation

Assume:

- A basic heap interface:
 - `malloc(size_t size)` allocates `size` or more bytes
 - `free(void *ptr)` frees corresponding chunk
 - Note that no size is provided
- Only **external fragmentation**
 - Allocators also have **internal fragmentation**
 - Unused space in chunks bigger than requested
- No memory relocation
 - No **compaction** of free space

Splitting



Splitting

`malloc(2048)`



Splitting

`malloc(2048)`



Splitting

```
malloc(2048)=15KB
```



Splitting



Coalescing

free (15KB)

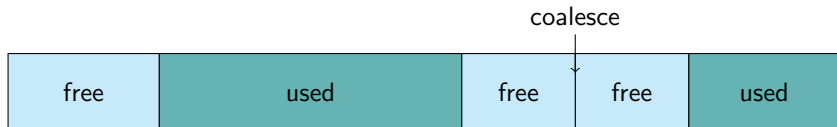


Coalescing

free (15KB)



Coalescing

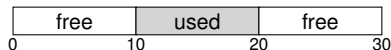


Coalescing

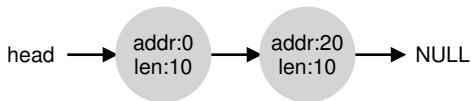


Splitting

- Find a free chunk to satisfy request and split it into two
 - Assume the following 30-byte heap:



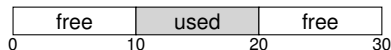
- Its free list is:



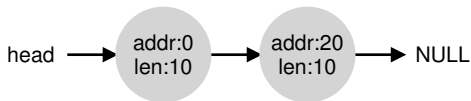
- After a 1-byte request:

Splitting

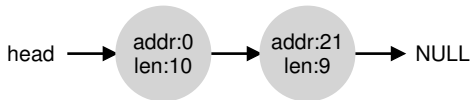
- Find a free chunk to satisfy request and split it into two
 - Assume the following 30-byte heap:



- Its free list is:

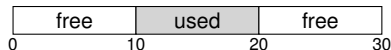


- After a 1-byte request:



Coalescing

- Coalesce free space when memory is freed
 - i.e., merge contiguous free chunk
 - Consider our previous heap:



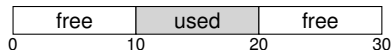
- After a call to `free(10)`:



- With coalescing:

Coalescing

- Coalesce free space when memory is freed
 - i.e., merge contiguous free chunk
 - Consider our previous heap:



- After a call to `free(10)`:


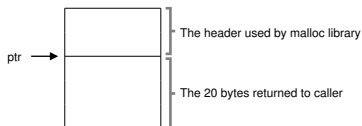


- With coalescing:



Tracking The Size

- Interface to `free(void *ptr)` does not provide size
- Store extra information in a **header** block
 - Usually just before chunk of memory
 - Magic number for integrity checking

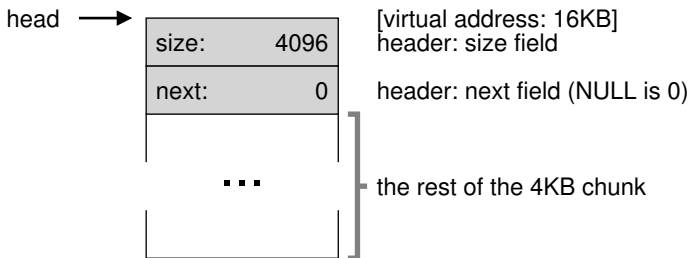


Free List

- Need to implement the **free list** itself
 - Can't call `malloc()` - we are implementing it!
 - Need to embed the list inside the free space itself


Free List

- Need to implement the **free list** itself
 - Can't call `malloc()` - we are implementing it!
 - Need to embed the list inside the free space itself
- Example: manage 4096-byte chunk
 - Maintain `size` and `next` for each node:



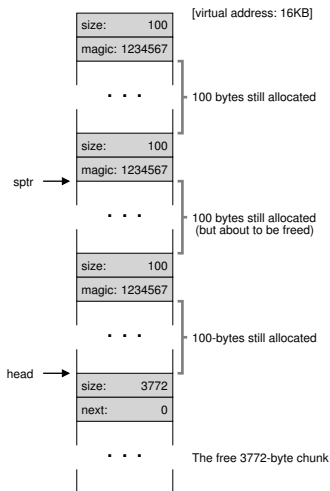
Free List

- 100 bytes are requested: **split** chunk



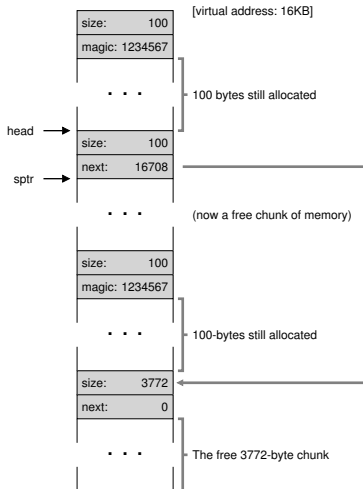
Free List

- Three allocated regions of 100 bytes:



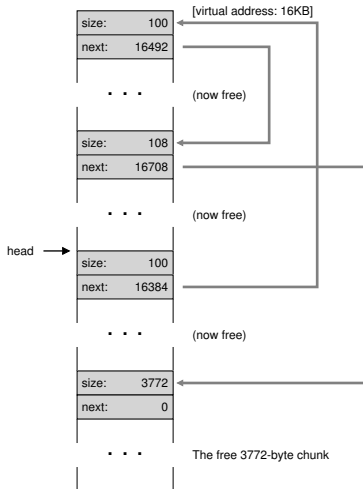
Free List

- After `free(16500)`:
 - Region start: 16384, previous chunk: 108, current header: 8



Free List

- Free last two chunks: fragmentation!
 - Need to **coalesce** the list



Managing Free Space

- **Best Fit**

- Return smallest chunk that's as big or bigger than requested size
- Exhaustive search: heavy performance penalty

- **Worst Fit**

- Opposite of best fit: return largest chunk
- Still requires full search, bad performance, excess fragmentation

Managing Free Space

- **First Fit**

- Return first block that is big enough
- Speed advantage, but pollutes beginning of list

- **Next Fit**

- As first fit, but start where stopped previously
- Spreads the searches throughout the list

Managing Free Space

- Examples: allocation request size 15



- Best-fit:



- Worst-fit:



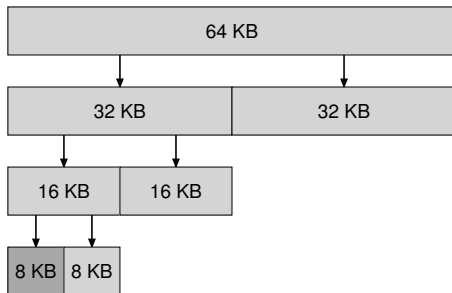
- First-fit: same as worst-fit, but faster

Segregated Lists

- Keep lists for fixed-size objects
- General memory allocator for the rest
- **Slab allocator** allocates **object caches**
 - For common kernel objects (locks, file-system inodes, etc.)
 - When a cache is running low: request **slab** of memory from general allocator

Buddy Allocation

- Make coalescing simple: **binary buddy allocator**
- Divide free space by two until a block is found
 - Further split into two is too small
 - Suffers from **internal fragmentation**
 - Easy to coalesce:
 - Recursively up the tree
 - Buddy address differs by a single bit



Summary

- Virtualize RAM into process **address space**
- Address translation:
 - Dynamic relocation (**base and bounds**)
 - Segmentation
 - **Coarse-grained**: just a few segments
 - **Fine-grained**: large number of smaller segments
- **External fragmentation**
 - Free memory fragments into small parts
 - Splitting & coalescing
- **Compaction**
 - Stop processes, copy data to contiguous region