

# Synchronization Primitives (ch. 30+31+32)

Operating Systems

Based on: Three Easy Pieces by Arpaci-Dusseau

Moshe Sulamy

Tel-Aviv Academic College

# Condition Variables

- Many cases a thread wishes to wait until a certain **condition**
- e.g., waiting for another thread to complete
  - Often called a `join()`
- Shared variable: works, but hugely inefficient

How should a thread wait for a condition?

# Condition Variables

- **Waiting** on a condition
  - Thread puts itself in a queue until some state of execution
- **Signaling** on a condition
  - Some other thread can wake waiting thread

# Condition Variables

- **Waiting** on a condition
  - Thread puts itself in a queue until some state of execution
- **Signaling** on a condition
  - Some other thread can wake waiting thread
- Name is a bit misleading
  - More of a queue
  - We are responsible for the actual “condition”

# Definitions and Routines

- Declare a condition variable:

```
pthread_cond_t cv;
```

- Operations:

```
// wait:  
pthread_cond_wait(pthread_cond_t *c, pthread_mutex_t *m);  
// signal:  
pthread_cond_signal(pthread_cond_t *c);
```

- Wait call takes **mutex** as a parameter
  - Caller must be its **owner** (have it locked)
  - Releases the lock, puts caller to sleep
  - On wake up, re-acquires lock and returns

# Parent Waiting For Child

- Two threads:
  - **Parent:**
    - Creates child thread
    - Waits on CV until child completes
  - **Child:**
    - Prints a message (“child”)
    - Wakes parent by signaling on CV

# Parent Waiting For Child

```
1 pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
2 pthread_cond_t c = PTHREAD_COND_INITIALIZER;
3
4 void* child(void* arg) {
5     printf("child\n");
6     thr_exit();
7     return NULL;
8 }
9 int main(void) {
10    printf("parent: begin\n");
11    pthread_t p;
12    pthread_create(&p, NULL, child, NULL);
13    thr_join();
14    printf("parent: end\n");
15    return 0;
16 }
```

# Parent Waiting For Child

```
1 void thr_exit() {  
2     pthread_cond_signal(&c);  
3 }  
4 void thr_join() {  
5     pthread_mutex_lock(&m);  
6     pthread_cond_wait(&c, &m);  
7     pthread_mutex_unlock(&m);  
8 }
```

- Why might this code fail?



# Parent Waiting For Child

```
1 void thr_exit() {  
2     pthread_cond_signal(&c);  
3 }  
4 void thr_join() {  
5     pthread_mutex_lock(&m);  
6     pthread_cond_wait(&c, &m);  
7     pthread_mutex_unlock(&m);  
8 }
```

- Why might this code fail?
  - Child runs immediately
  - Will signal, but no thread asleep on CV
  - Parent runs, calls wait and gets stuck
  - Solution?

# Parent Waiting For Child

```
1 void thr_exit() {  
2     pthread_cond_signal(&c);  
3 }  
4 void thr_join() {  
5     pthread_mutex_lock(&m);  
6     pthread_cond_wait(&c, &m);  
7     pthread_mutex_unlock(&m);  
8 }
```

- Why might this code fail?
  - Child runs immediately
  - Will signal, but no thread asleep on CV
  - Parent runs, calls wait and gets stuck
  - Solution? use done variable

# Parent Waiting For Child

```
1  int done = 0;
2  void thr_exit() {
3      done = 1;
4      pthread_cond_signal(&c);
5  }
6  void thr_join() {
7      pthread_mutex_lock(&m);
8      if (done == 0)
9          pthread_cond_wait(&c, &m);
10     pthread_mutex_unlock(&m);
11 }
```

- Why might this code fail?

# Parent Waiting For Child

```
1  int done = 0;
2  void thr_exit() {
3      done = 1;
4      pthread_cond_signal(&c);
5  }
6  void thr_join() {
7      pthread_mutex_lock(&m);
8      if (done == 0)
9          pthread_cond_wait(&c, &m);
10     pthread_mutex_unlock(&m);
11 }
```

- Why might this code fail?
  - Parent calls join, sees done=0
  - Interrupted just before wait, context switch to child
  - Child sets done, signal is lost, parent is stuck again
  - Solution?

# Parent Waiting For Child

```
1  int done = 0;
2  void thr_exit() {
3      done = 1;
4      pthread_cond_signal(&c);
5  }
6  void thr_join() {
7      pthread_mutex_lock(&m);
8      if (done == 0)
9          pthread_cond_wait(&c, &m);
10     pthread_mutex_unlock(&m);
11 }
```

- Why might this code fail?
  - Parent calls join, sees done=0
  - Interrupted just before wait, context switch to child
  - Child sets done, signal is lost, parent is stuck again
  - Solution? hold lock while signaling

# Parent Waiting For Child

```
1  int done = 0;
2  void thr_exit() {
3      pthread_mutex_lock(&m);
4      done = 1;
5      pthread_cond_signal(&c);
6      pthread_mutex_unlock(&m);
7  }
8  void thr_join() {
9      pthread_mutex_lock(&m);
10     while (done == 0)
11         pthread_cond_wait(&c, &m);
12     pthread_mutex_unlock(&m);
13 }
```

- Additionally, check variable in a loop
  - Condition variable may signal unexpectedly
  - Also crucial for more than 2 threads

# Covering Conditions

- Memory allocator implementation
- Assume zero bytes are free:
  - Thread A calls `allocate(100)`
  - Thread B calls `allocate(10)`
  - Both A and B wait on the condition
  - Thread C calls `free(50)`

# Covering Conditions

- Memory allocator implementation
- Assume zero bytes are free:
  - Thread A calls `allocate(100)`
  - Thread B calls `allocate(10)`
  - Both A and B wait on the condition
  - Thread C calls `free(50)`
    - Which waiting thread wakes up?



# Covering Conditions

- Memory allocator implementation
- Assume zero bytes are free:
  - Thread A calls `allocate(100)`
  - Thread B calls `allocate(10)`
  - Both A and B wait on the condition
  - Thread C calls `free(50)`
    - Which waiting thread wakes up?
- Solution: wake up all waiting threads
  - Use `pthread_cond_broadcast()`
  - Performance cost: too many threads might be woken

# Producer / Consumer

- Also: **bounded buffer** problem
- **Producer**
  - Produces data items
  - Wishes to place items in a buffer
- **Consumer**
  - Grabs items out of the buffer
  - Consumes items in some way
- e.g., web server consumes HTTP requests (work queue)

# Producer / Consumer

- Also used in pipes:
  - `grep foo file.txt | wc -l`
  - `grep` - output lines from `file.txt` containing `foo`
  - `wc -l` - output number of lines from input
  - Shell redirects `grep` standard output to a **pipe**
    - Created by the `pipe` system call
  - Other end connected to standard input of `wc`
  - `grep` - producer, `wc` - consumer
- In-kernel bounded buffer

# Producer / Consumer

```
1 int buffer;  
2 int count = 0; // initially, empty  
3  
4 void put(int value) {  
5     assert(count == 0);  
6     count = 1;  
7     buffer = value;  
8 }  
9 int get() {  
10    assert(count == 1);  
11    count = 0;  
12    return buffer;  
13 }
```

# First Attempt

- What are the two problems?

```
1 pthread_cond_t cond;
2 pthread_mutex_t mutex;
3
4 void produce(int i) {
5     pthread_mutex_lock(&mutex);
6     if (count == 1)
7         pthread_cond_wait(&cond, &mutex);
8     put(i);
9     pthread_cond_signal(&cond);
10    pthread_mutex_unlock(&mutex);
11 }
12 int consume() {
13     pthread_mutex_lock(&mutex);
14     if (count == 0)
15         pthread_cond_wait(&cond, &mutex);
16     int tmp = get();
17     pthread_cond_signal(&cond);
18     pthread_mutex_unlock(&mutex);
19     return tmp;
20 }
```

# First Attempt

- More than one consumer / producer

# First Attempt

- More than one consumer / producer
- Two consumers:
  - No while on CV in consume()
  - Can consume when empty!
  - (and the same for produce())

```
1 int consume() {  
2     pthread_mutex_lock(&mutex);  
3     while (count == 0)  
4         pthread_cond_wait(&cond, &mutex);  
5     int tmp = get();  
6     pthread_cond_signal(&cond);  
7     pthread_mutex_unlock(&mutex);  
8     return tmp;  
9 }
```

# First Attempt

buffer=0

Producer

Consumer 1

Consumer 2



# First Attempt

buffer=0

Producer

Consumer 1

consume

Consumer 2

# First Attempt

buffer=0

Producer

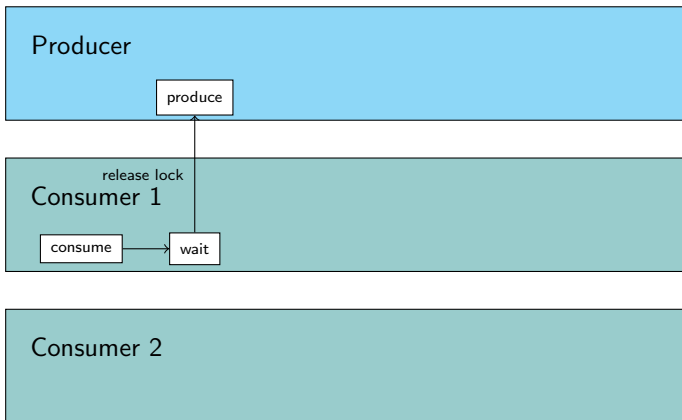
Consumer 1



Consumer 2

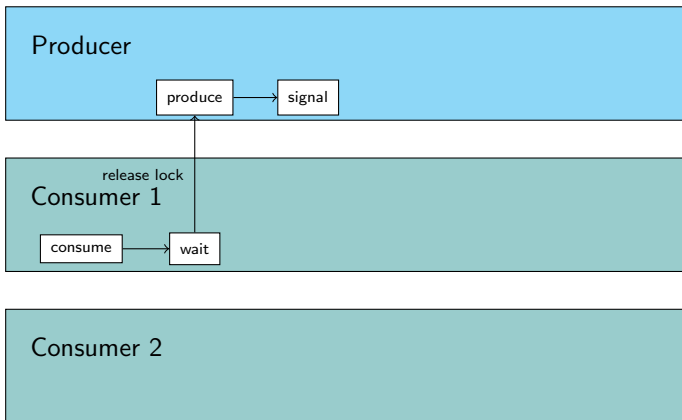
# First Attempt

buffer=1



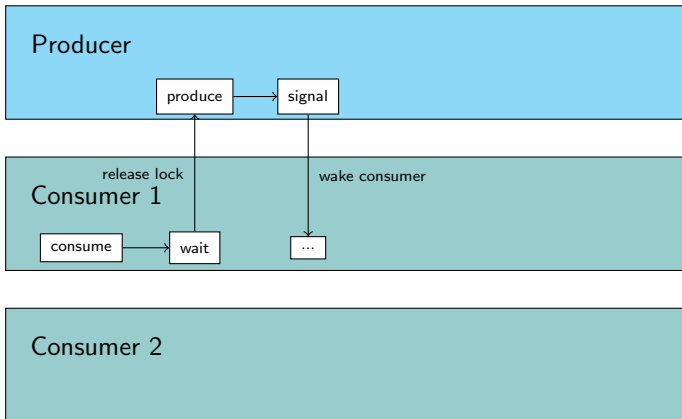
# First Attempt

buffer=1



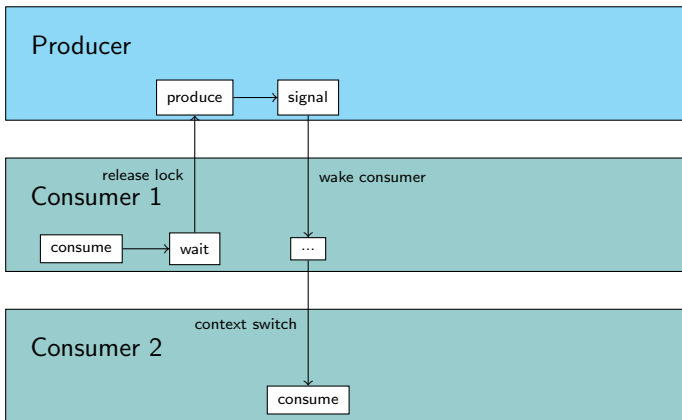
# First Attempt

buffer=1



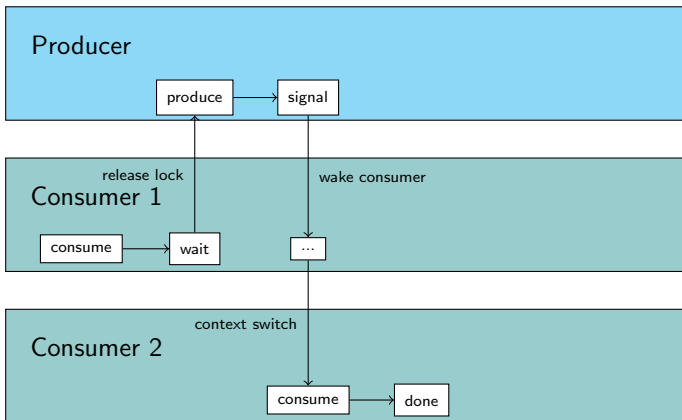
# First Attempt

buffer=1



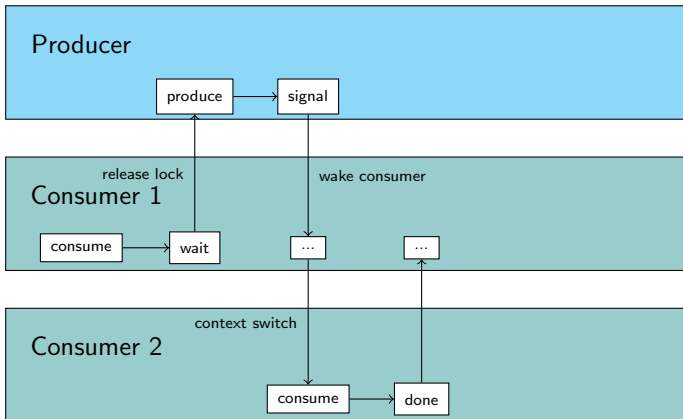
# First Attempt

buffer=0



# First Attempt

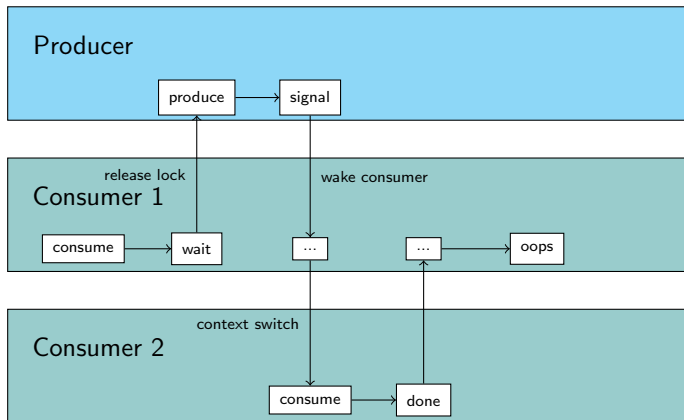
buffer=0





# First Attempt

buffer=0



# First Attempt

- Only one condition variable!
  - Consumer might wake another consumer
  - Producer might wake another producer
- Solution?

# First Attempt

- Only one condition variable!
  - Consumer might wake another consumer
  - Producer might wake another producer
- Solution? use **two** condition variables
  - Producer threads wait on `empty`
  - Consumer threads wait on `full`

# Semaphores

- Another **synchronization primitive**
  - Can use **semaphores** as both locks and condition variables
- A **semaphore**
  - Object with an integer value
  - `sem_wait()`, `sem_post()`

# Semaphores

- Initialization:

```
#include <semaphore.h>
sem_t s;
sem_init(&s, 0, 1);
```

- Initialized to the value 1 (third argument)
  - Second argument: 0 to share between threads (vs. processes)
  - We will only use 0

# Semaphores

- `sem_wait()`:
  - Decrement semaphore value by one
  - Wait if value is negative
- `sem_post()`:
  - Increment semaphore value by one
  - Wake one waiting thread
- Both operations are performed **atomically**

# Binary Semaphores

- How can we use a semaphore as a lock?

```
1 sem_init(&m, 0, X); // what should X be?  
2 ...  
3 sem_wait(&m);  
4 // critical section  
5 sem_post(&m);
```

- What should **X** be?

# Binary Semaphores

- How can we use a semaphore as a lock?

```
1 sem_init(&m, 0, X); // what should X be?  
2 ...  
3 sem_wait(&m);  
4 // critical section  
5 sem_post(&m);
```

- What should **X** be?
  - $X = 1$
  - **Binary semaphore**



# Ordering Primitive

- Similar to **condition variables**
- How can we use a semaphore to wait for an event?

```
1 sem_t s;  
2  
3 void* child(void* arg) {  
4     printf("child\n");  
5     sem_post(&s); // child is done  
6     return NULL;  
7 }  
8 int main() {  
9     sem_init(&s, 0, X); // what should X be?  
10    printf("parent: begin\n");  
11    pthread_t c;  
12    pthread_create(&c, NULL, child, NULL);  
13    sem_wait(&s); // wait for child  
14    printf("parent: end\n");  
15    return 0;  
16 }
```

- What should **X** be?

# Ordering Primitive

- Similar to **condition variables**
- How can we use a semaphore to wait for an event?

```
1 sem_t s;  
2  
3 void* child(void* arg) {  
4     printf("child\n");  
5     sem_post(&s); // child is done  
6     return NULL;  
7 }  
8 int main() {  
9     sem_init(&s, 0, X); // what should X be?  
10    printf("parent: begin\n");  
11    pthread_t c;  
12    pthread_create(&c, NULL, child, NULL);  
13    sem_wait(&s); // wait for child  
14    printf("parent: end\n");  
15    return 0;  
16 }
```

- What should **X** be?
  - X=0

# Semaphores

- What value should a semaphore be initialized to?
  - General rules: number of resources
  - Lock: 1  $\rightarrow$  can be locked after initialization
  - Ordering: 0  $\rightarrow$  nothing to give away at the start

# Producer / Consumer

- How can we implement a **bounded buffer** with semaphores?

```
1  int buffer[MAX];
2  int fill = 0;
3  int use  = 0;
4
5  void put(int value) {
6      buffer[fill] = value;
7      fill = (fill + 1) % MAX;
8  }
9  int get() {
10     int tmp = buffer[use];
11     use = (use + 1) % MAX;
12     return tmp;
13 }
```

# First Attempt

- What is the problem?

```
1 sem_t empty; // initialized to MAX
2 sem_t full;  // initialized to 0
3
4 void produce(int value) {
5     sem_wait(&empty);
6     put(value);
7     sem_post(&full);
8 }
9 int consume() {
10    sem_wait(&full);
11    int tmp = get();
12    sem_post(&empty);
13    return tmp;
14 }
```

# First Attempt

Producer 1

Producer 2

# First Attempt


Producer 1

```
put(5)
```

Producer 2

# First Attempt

Producer 1




Producer 2




# First Attempt

Producer 1

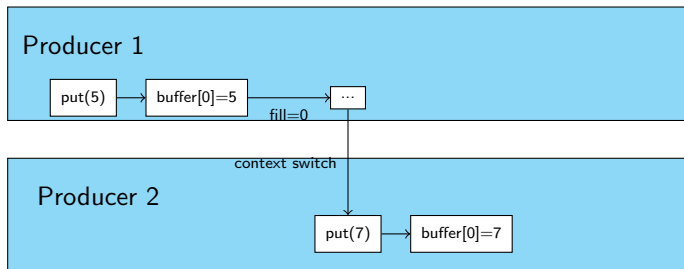


Producer 2

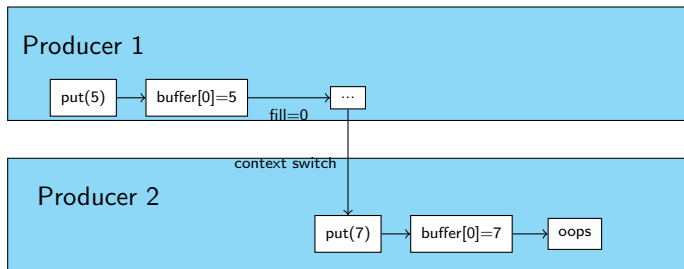
# First Attempt



# First Attempt



# First Attempt



# With Mutual Exclusion

- What is the problem?

```
1  sem_t mutex; // binary semaphore
2  sem_t empty; // initialized to MAX
3  sem_t full;  // initialized to 0
4
5  void produce(int value) {
6      sem_wait(&mutex);
7      sem_wait(&empty);
8      put(value);
9      sem_post(&full);
10     sem_post(&mutex);
11 }
12 int consume() {
13     sem_wait(&mutex);
14     sem_wait(&full);
15     int tmp = get();
16     sem_post(&empty);
17     sem_post(&mutex);
18     return tmp;
19 }
```

# With Mutual Exclusion

- What is the problem? **deadlock!**
  - Producer waits on empty, holds mutex, consumer can't consume

```
1 sem_t mutex; // binary semaphore
2 sem_t empty; // initialized to MAX
3 sem_t full; // initialized to 0
4
5 void produce(int value) {
6     sem_wait(&mutex);
7     sem_wait(&empty);
8     put(value);
9     sem_post(&full);
10    sem_post(&mutex);
11 }
12 int consume() {
13     sem_wait(&mutex);
14     sem_wait(&full);
15     int tmp = get();
16     sem_post(&empty);
17     sem_post(&mutex);
18     return tmp;
19 }
```

# With Mutual Exclusion

- Solution: use mutex around the critical section

```
1 sem_t mutex; // binary semaphore
2 sem_t empty; // initialized to MAX
3 sem_t full; // initialized to 0
4
5 void produce(int value) {
6     sem_wait(&empty);
7     sem_wait(&mutex);
8     put(value);
9     sem_post(&mutex);
10    sem_post(&full);
11 }
12 int consume() {
13     sem_wait(&full);
14     sem_wait(&mutex);
15     int tmp = get();
16     sem_post(&mutex);
17     sem_post(&empty);
18     return tmp;
19 }
```

# Reader-Writer Locks

- More flexible locking primitive
  - e.g., concurrent operations: inserts and lookups
  - Insert changes state → traditional critical section
  - Lookup reads data structure → many at once (if no insert)
- **Reader-writer lock**
  - Four operations: acquire/release read/write lock



# Reader-Writer Locks

- A **single writer** can acquire the lock
- Once a reader acquires a **read lock**:
  - **More readers** are allowed to acquire the read lock
  - A writer waits until all readers are finished

# Reader-Writer Locks

- A **single writer** can acquire the lock
- Once a reader acquires a **read lock**:
  - **More readers** are allowed to acquire the read lock
  - A writer waits until all readers are finished

```
1 typedef struct _rwlock_t {
2     sem_t lock;           // binary semaphore
3     sem_t writelock;      // used to allow ONE writer
4     int readers;          // count of readers in CS
5 } rwlock_t;
6
7 void rwlock_init(rwlock_t* rw) {
8     rw->readers = 0;
9     sem_init(&rw->lock, 0, 1);
10    sem_init(&rw->writelock, 0, 1);
11 }
```

# Reader-Writer Locks

```
1 void rwlock_acquire_writelock(rwlock_t* rw) {
2     sem_wait(&rw->writelock);
3 }
4 void rwlock_release_writelock(rwlock_t* rw) {
5     sem_post(&rw->writelock);
6 }
7
8 void rwlock_acquire_readlock(rwlock_t* rw) {
9     sem_wait(&rw->lock);      // CS for readers
10    rw->readers++;
11    if (rw->readers == 1)
12        sem_wait(&rw->writelock); // first reader grabs writelock
13    sem_post(&rw->lock);
14 }
15 void rwlock_release_readlock(rwlock_t* rw) {
16     sem_wait(&rw->lock);      // CS for readers
17    rw->readers--;
18    if (rw->readers == 0)
19        sem_post(&rw->writelock); // last reader releases writelock
20    sem_post(&rw->lock);
21 }
```

# Reader-Writer Locks


- What is the problem?

# Reader-Writer Locks

- What is the problem? **fairness**
  - Easy to **starve** writer
  - How to prevent readers from starving writers?

# Dining Philosophers

- Five philosophers around a table
  - Single fork between each pair
  - Philosophers **think** and **eat**
  - Two forks to eat (left and right)



# Dining Philosophers

- As code:

```
1 while (1) {  
2     think();  
3     getforks();  
4     eat();  
5     putforks();  
6 }
```

```
1 int left(int p) {  
2     return p;  
3 }  
4 int right(int p) {  
5     return (p - 1) % 5;  
6 }
```

# Dining Philosophers

- Use a semaphore for each fork:

```
1 void get_forks(int p) {  
2     sem_wait(&forks[left(p)]);  
3     sem_wait(&forks[right(p)]);  
4 }  
5 void put_forks(int p) {  
6     sem_post(&forks[left(p)]);  
7     sem_post(&forks[right(p)]);  
8 }
```

- The problem?



# Dining Philosophers

- Use a semaphore for each fork:

```
1 void get_forks(int p) {  
2     sem_wait(&forks[left(p)]);  
3     sem_wait(&forks[right(p)]);  
4 }  
5 void put_forks(int p) {  
6     sem_post(&forks[left(p)]);  
7     sem_post(&forks[right(p)]);  
8 }
```

- The problem? **deadlock!**
  - Each philosopher grabs fork on their left
  - All waiting for their right

# Dining Philosophers

- Solution: break the dependency

```
1 void get_forks(int p) {  
2     if (p == 4) {  
3         sem_wait(&forks[right(p)]);  
4         sem_wait(&forks[left(p)]);  
5     }  
6     else {  
7         sem_wait(&forks[left(p)]);  
8         sem_wait(&forks[right(p)]);  
9     }  
10 }
```

# Thread Throttling

- For example: hundreds of threads work in parallel
- Section of code allocates a lot of memory
  - All threads at the same time → exceeds physical memory
  - Machine will start thrashing (swapping to and from the disk)
- Solution?

# Thread Throttling

- For example: hundreds of threads work in parallel
- Section of code allocates a lot of memory
  - All threads at the same time → exceeds physical memory
  - Machine will start thrashing (swapping to and from the disk)
- Solution? **semaphore!**
  - Initialized to max threads we wish to enter code section
  - Surrounds code section, limits concurrent threads in it

# Implementing Semaphores

- Doesn't maintain invariant: negative value  $\rightarrow$  # waiting threads
  - Easier, matches the Linux implementation

```
1  typedef struct __zem_t {
2      int value;
3      pthread_cond_t cond;
4      pthread_mutex_t lock;
5  };
6
7  void zem_init(zem_t* s, int value) {
8      s->value = value;
9      pthread_cond_init(&s->cond);
10     pthread_mutex_init(&s->lock);
11 }
12 void zem_wait(zem_t* s) {
13     pthread_mutex_lock(&s->lock);
14     while (s->value <= 0)
15         pthread_cond_wait(&s->cond, &s->lock);
16     s->value--;
17     pthread_mutex_unlock(&s->lock);
18 }
19 void zem_post(zem_t* s) {
20     pthread_mutex_lock(&s->lock);
21     s->value++;
22     pthread_cond_signal(&s->cond);
23     pthread_mutex_unlock(&s->lock);
24 }
```

# Summary

- **Condition variables**

- Thread waits until a certain condition
- `wait()`, `signal()`
- **Hold lock** while signaling
- Check value **in a loop**

- **Semaphore**

- Integer value
- Decrement on acquire, wait if negative, increment on release

- **Read-write lock**

- **Single writer** or **multiple readers**

- Dining philosophers

- Think and eat

- Producer / consumer (bounded buffer)